

Krishnakumar Ramesan

Hayward, CA | (256) 652-8738 | ramesankrishnakumar@gmail.com | LinkedIn | GitHub

May 29, 2026

Paramount Streaming Hiring Team
Paramount Global
Los Angeles, CA

Dear Hiring Team,

I'm applying for the Lead Software Engineer, Backend Services role on Paramount Streaming's team. The role centers on leading API development for a real-time personalization engine that decides the optimal experience for each viewer, and that maps closely to what I've spent the last few years building. At Intuit I designed and shipped a production multi-agent AI system for QuickBooks onboarding on LangGraph and GPT 4.1, taking it from hackathon prototype to serving 50% of first-time trial customers as tech lead. At its core is a config-driven intent orchestrator that classifies each request in real time and routes it across seven domain agents, which is the same real-time decisioning shape your personalization engine calls for.

Low-latency under load is where I focus. For that same system I replaced a screenshot-plus-vision-LLM pipeline of roughly 9 seconds per turn with a 30 ms DOM parse, and on the data side I cut a bulk catalog-matching API from a 30-second timeout to 785 ms, a 97% reduction, by profiling query patterns and removing redundant round-trips and missing indexes. I build APIs in Kotlin and Java on Spring Boot, and I've shipped GraphQL services with Netflix DGS, so your Java and contract-first stack is familiar ground.

Your role leans heavily on event-driven architecture, and I've owned that end to end. At Intuit I built a fault-tolerant fulfillment subsystem ingesting events from four sales channels, hundreds of thousands per day, designed for idempotent consumers and ordered delivery and scaled to 75 TPS through partitioning, sharding, and load-based autoscaling on Kafka. Earlier at TriNet I built a core SSO microservice serving roughly 400K customers across 20 provider integrations and an ACL authorization service consumed by 20 microservices, cutting p90 latency with a two-tier cache, so high-scale APIs and clean cross-team contracts are routine for me.

On the AI-in-the-SDLC point in your posting, I use tools like Claude Code daily for development, testing, and debugging, and I've also built the AI systems themselves, so I bring both the practitioner and the builder view. I'd welcome the chance to lead API work that shapes how millions of viewers experience Paramount+ and Pluto TV.

Sincerely,
Krishnakumar Ramesan